

BEYOND MEAN FOILS: AUDITING WORST-FOIL SPECIFICITY IN FROZEN CLIP REGION EXPLANATIONS

Kaixin Liu^{1,*}, Zhipeng Ye^{1,*†}, Feng Jiang¹, Zhenghao Wang¹, Qihang Wu¹

¹Taizhou Institute of Science and Technology, Nanjing University of Science and Technology,
Taizhou 225300, Jiangsu, China

24107880127@nustti.edu.cn, zhipengye@nustti.edu.cn, jf@nustti.edu.cn,
wangzhenghao2002@outlook.com, 24107880128@nustti.edu.cn

*Equal contribution. †Corresponding author: Zhipeng Ye.

ABSTRACT

A region can overlap a target object yet contribute more to another class. We test regions selected by Cluster-based Concept Importance (CCI) in frozen CLIP. Across COCO and VOC with two checkpoints, 41.08–64.78% of regions that pass overlap and mean-contrast checks fail against the strongest competing class. Removing competitors annotated in the image leaves 39.69–63.64% failing. We then test all eight candidate regions per image. An alternative passes the test for 6.25–7.84% of failures on COCO and 27.40–31.15% on VOC. Requiring it to preserve the original target-score drop within $\epsilon = .02$ reduces these rates to 0.16–0.98%. Available regions and target-drop tolerance constrain repair; relaxing the tolerance increases repair opportunities.

Index Terms— CLIP, explanation auditing, class specificity, constrained region selection

1 INTRODUCTION

Cluster-based Concept Importance (CCI) [2] explains CLIP [1] by grouping patches into regions and measuring how masking them changes image–text scores. We examine CCI-top1, the region whose removal most reduces the target score.

A region may contribute more to another class than to the target, even when its target drop exceeds the average competing drop. We test the strongest competitor, or worst foil, and ask whether another region can pass this test while preserving the target drop.

Class-contrastive explanations [3] can choose a competitor by its score in the original image. Our audit instead chooses the class whose score drops most when a region is masked. Related work covers corpus foils [4], class-sensitive saliency [5], CLIP attribution [6, 7, 8], class contrast [9], and caption foils [10]. We control for competing objects and foil count, then examine all regions to locate constraints on repair.

2 AUDIT AND REPAIR OPPORTUNITY

Figure 1 shows candidate construction (a), one-prompt selection (b), and three-prompt evaluation of the fixed region (c).

For image x , let t be its target and c any class. The candidate set $\mathcal{R}(x)$ contains $K = 8$ patch regions R . Checkpoint

m gives pre-softmax score $s_m(x, c)$, including its learned logit scale. Masking attention to R gives $\mathcal{I}(x, R)$. The response $d_R(c)$ is the original score minus the masked score; positive values mean a reduction:

$$d_R(c) = s_m(x, c) - s_m(\mathcal{I}(x, R), c). \quad (1)$$

Superscript sel denotes “a photo of a {category}”. CCI-top1 chooses R_{CCI} , the region with the largest target drop; arg max returns the maximizing region:

$$R_{CCI} = \arg \max_{R \in \mathcal{R}(x)} d_R^{\text{sel}}(t). \quad (2)$$

Evaluating class specificity. We evaluate the selected region with “a picture of the {category}”, “an image containing a {category}”, and “the {category}”. Index the $H = 3$ prompts by h . Their unit-norm text embeddings are $e_h(c)$ and their drops are $d_{R,h}(c)$. The mean embedding has length $n_c = \|H^{-1} \sum_h e_h(c)\|_2$, where $\|\cdot\|_2$ is Euclidean length. Dividing the mean drop by this length gives the drop for a normalized mean embedding: $\bar{d}_c(R) = H^{-1} \sum_h d_{R,h}(c)/n_c$.

A foil is a non-target class, and $\mathcal{F}_t$ is the set of all such classes. With f indexing foils and R implicit in the drops, margin $P(R, t)$ subtracts the largest foil drop from the target drop. Negative values mean that a foil loses more score:

$$P(R, t) = \bar{d}_t - \max_{f \in \mathcal{F}_t} \bar{d}_f. \quad (3)$$

We call $P < 0$ a failure and a transition to $P \geq 0$ a sign repair. Target-response tables use the raw mean drop $H^{-1} \sum_h d_{R,h}(t)$.

For baseline CCI-top1, A contains images passing both checks: bbox precision (the fraction of patch centers inside target boxes) is at least .5, and the target drop exceeds the mean foil drop. For image i with target t_i , write $P_i(R) = P(R, t_i)$; P_i denotes its baseline margin. Failures form $F = \{i : P_i < 0\}$, so the screened failures are $B = A \cap F$. These image sets stay fixed. The subset with positive target contribution is $A^+ = A \cap \{\bar{d}_t > 0\}$.

Selecting an alternative region. Worst-foil (WF) reranking allows a target-drop loss of at most $\epsilon \geq 0$ relative to CCI-top1 (Figure 1(b)). Its selection margin M^{sel} subtracts the largest




Figure 1: Selection and held-out audit. A region is fixed using selection-stage responses before three-template evaluation of specificity, target response, and locality. Schematic responses illustrate the selection rule.

foil drop from the target drop under the selection prompt:

$$M^{\text{sel}}(R, t) = d_R^{\text{sel}}(t) - \max_{f \in \mathcal{F}_t} d_R^{\text{sel}}(f), \quad (4)$$

The feasible set $\mathcal{R}_\epsilon$ contains candidates within the allowed target-drop loss:

$$\mathcal{R}_\epsilon = \{R \in \mathcal{R}(x) : d_R^{\text{sel}}(t) \geq d_{R_{\text{CCI}}}^{\text{sel}}(t) - \epsilon\}, \quad (5)$$

WF chooses R_{WF} to maximize that margin within the feasible set:

$$R_{\text{WF}} = \arg \max_{R \in \mathcal{R}_\epsilon} M^{\text{sel}}(R, t). \quad (6)$$

We use $\epsilon = .02$ logit units and sweep $.01/.02/.05/.10/.20$.

Locating the constraints on repair. An oracle tests whether any allowed alternative can repair a failure. It counts images in B with a feasible region with a nonnegative held-out margin. Denote this count by $N_O(B, \epsilon)$ and image i 's feasible set by $\mathcal{R}_\epsilon(i)$. The indicator $\mathbf{1}\{\cdot\}$ is one when its condition holds and zero otherwise:

$$N_O(B, \epsilon) = \sum_{i \in B} \mathbf{1}\left\{ \max_{R \in \mathcal{R}_\epsilon(i)} P_i(R) \geq 0 \right\}. \quad (7)$$

The unrestricted oracle removes the budget constraint and searches all K candidates. Each image in B falls into one of four groups: no candidate has $P \geq 0$ (C_0); a candidate passes but none meets the budget (C_1); a feasible candidate passes but WF misses it (C_2); or WF repairs the failure (r). These groups partition B , whose size is $|B|$, so $|B| = C_0 + C_1 + C_2 + r$. The oracle counts both feasible groups: $N_O = C_2 + r$. The joint condition J additionally requires overlap and mean-foil checks and $\bar{d}_t > 0$; j counts WF repairs satisfying it within fixed B . WF selects using its selection prompt; only the oracle uses held-out scores.

Controlling for competing objects. A foil may name another object in the image. To measure its contribution, remove all annotated-present classes from image i 's foil set, leaving $\mathcal{F}_i^{\text{abs}}$. The protocol holds R_{CCI} and full-foil screen A fixed; the failure event is $E_i^{\text{abs}} = \{\bar{d}_t - \max_{f \in \mathcal{F}_i^{\text{abs}}} \bar{d}_f < 0\}$. Re-

moving any foils can reduce failure, so we compare this exclusion with randomly retaining the same number of foils. For N_i full foils, $k_i = |\mathcal{F}_i^{\text{abs}}|$, and m_i foils with $\bar{d}_f > \bar{d}_t$, let q_i be the failure probability of a uniform size- k_i subset. With $\binom{n}{k}$ counting size- k subsets, one minus the fraction containing no stronger foil gives

$$q_i = 1 - \frac{\binom{N_i - m_i}{k_i}}{\binom{N_i}{k_i}}. \quad (8)$$

The numerator is zero when $k_i > N_i - m_i$. To isolate the effect of annotation-based removal, Δ_A averages its failure indicator minus the random-subset probability over the $|A|$ baseline images: $\Delta_A = |A|^{-1} \sum_{i \in A} (\mathbf{1}\{E_i^{\text{abs}}\} - q_i)$. Negative Δ_A favors annotation-based removal. Retained foil sets are nonempty.

3 PROTOCOL AND REPRODUCIBILITY

We evaluate 27,708 COCO val2014 [11] and 1,943 VOC2007 test [12] images with CLIP B/16 and B/32 (16- and 32-pixel patches): 59,302 image-model records. The target is the largest instance by annotation area (COCO) or box area (VOC). Images require another category and 2–80% target area; COCO crowd annotations are excluded. Full foils are the other 79/19 classes.

K-means uses normalized projected patch features, $K = 8$, three starts, 50 iterations, and seed $1701 + i$; i indexes sorted IDs shuffled with seed 1701. Intervention $\mathcal{I}$ adds $-\infty$ to attention logits at selected patch-key columns before softmax in every block/head, retaining the class-token column. Ties use first argmax. Boxes follow the 224-pixel short-side resize and center crop [13].

4 RESULTS AND ANALYSIS

Table 1: Fixed baseline A : full-foil, annotation-absent, and exact cardinality-matched random failure rates (%). $100\Delta_A$ is absent minus random in percentage points, with paired 95% CI.

Setting	Full	Absent	Random	$100\Delta_A$
COCO B/16	64.37	63.17	64.09	-0.92 [-1.08,-0.77]
COCO B/32	64.78	63.64	64.50	-0.86 [-1.01,-0.71]
VOC B/16	42.27	40.94	41.46	-0.52 [-1.16,+0.06]
VOC B/32	41.08	39.69	40.24	-0.55 [-1.21,+0.04]

Removing annotated competitors leaves most failures.

Table 1 compares the same baseline images before and after removing classes annotated in the image. The set A contains 17,726, 17,848, 1,209, and 1,227 images in table order. Failure falls by only 1.14–1.39 percentage points: 96.63–98.24% of the failures in B remain. On COCO, exclusion helps more than removing the same number of classes at random. On VOC, the confidence intervals for that difference include zero. Annotated competing objects thus explain only a small part of the observed failure.

The pattern also holds when the target contribution is positive. In A^+ , failure after exclusion remains 61.38–62.62% on COCO and 39.26–40.05% on VOC. Even COCO images with exactly two annotated categories have failure rates of 63.39% and 64.95% for B/16 and B/32, respectively. The mismatch therefore persists beyond crowded scenes.

Table 2: Exact local decomposition of *screened failures* $B = A \cap F$ at $\epsilon = .02$. Counts in C_0, C_1, C_2, r are mutually exclusive; $C_2 + r$ is the feasible-oracle count.

Setting	$ B $	C_0	C_1	C_2	r
COCO B/16	11410	10697	695	0	18
COCO B/32	11562	10655	876	1	30
VOC B/16	511	371	135	0	5
VOC B/32	504	347	154	1	2

Few failures have an alternative within the budget. Table 2 shows where repair stops. For 92.16–93.75% of COCO failures and 68.85–72.60% of VOC failures, none of the eight candidates passes the worst-foil test (C_0). Figure 2 shows the remaining opportunity as blue circles: 6.25–7.84% on COCO and 27.40–31.15% on VOC. Imposing $\epsilon = .02$ removes 96.43–98.09% of these opportunities, leaving only 0.16–0.98% of B repairable (orange diamonds). WF misses at most one of these feasible repairs per setting (C_2). The main losses occur before the final ranking decision.




Figure 2: Budget-constrained sign-repair capacity among fixed screened failures B . Connected markers compare the same setting on a shared linear scale. Unrestricted: at least one of all eight candidates has $P \geq 0$; feasible: at least one also meets $\epsilon = .02$. Percentages use each setting’s own $|B|$ (Table 2) and are computed by exact enumeration.

Table 3: Budget endpoints (C: COCO; V: VOC). Multi: percentage of all eligible images with multiple feasible candidates. O_B, r_B : oracle/actual sign-repair percentages of screened failures B . $\Delta d_t, \Delta b$: all-image raw-target and bbox changes (bbox in percentage points).

Set	ϵ	Multi	O_B	r_B	Δd_t	Δb
C16	0.02	4.71	0.16	0.16	+0.0001	-0.12
C16	0.20	35.78	1.57	1.50	-0.0140	-1.63
C32	0.02	4.40	0.27	0.26	-0.0001	-0.08
C32	0.20	34.95	2.05	2.00	-0.0143	-0.60
V16	0.02	4.79	0.98	0.98	-0.0008	-0.34
V16	0.20	34.48	6.65	6.07	-0.0223	+0.08
V32	0.02	4.73	0.60	0.40	-0.0007	+0.02
V32	0.20	34.64	7.14	6.75	-0.0161	+1.07

A larger tolerance opens up more alternatives. At $\epsilon = .02$, about 95% of images have only one feasible candidate (Table 3). Increasing the tolerance to .20 gives multiple choices on about 35% of images. Oracle repair capacity rises to 1.57–2.05% of COCO failures and 6.65–7.14% of VOC failures. The oracle-repair percentage O_B and achieved WF percentage r_B are computed as $O_B = 100N_O/|B|$ and $r_B = 100r/|B|$.

The extra choice changes what the selected region captures. The table’s Δd_t and Δb are WF minus CCI changes averaged over all eligible images; bbox changes are in percentage points. At .20, raw target drops decrease in all four settings, while overlap decreases on COCO and increases on VOC. More repair opportunities therefore do not imply a uniform improvement in target response or localization.

Table 4: Normalization sensitivity on fixed earlier-run A_{old} and earlier-run CCI regions. Raw/Norm: failure rates (%); Flip: percentage whose failure status differs. A_{old} denotes earlier-run screen-passers.

Setting	$ A_{\text{old}} $	Raw	Norm	Flip
COCO B/16	17,726	64.318	64.369	0.107
COCO B/32	17,847	64.762	64.778	0.073
VOC B/16	1,209	42.184	42.266	0.083
VOC B/32	1,227	41.157	41.157	0.000

Normalization rarely changes the outcome. Table 4 isolates normalization by fixing earlier-run regions and images; these records are not pooled with rerun results. Failure rates differ by at most 0.083 percentage points, and at most 0.107% of images switch between passing and failing. The worst foil is unchanged on 99.42–99.84% of images. The high failure rate therefore persists under both score definitions.

Table 5: Actual sign repairs *within fixed baseline B*, $\epsilon = .02$. r counts repairs; j counts those also satisfying J . WF–CCI means with repair-subset 95% CIs. Δd_t is raw held-out target-score change; only bbox change is in percentage points.

Setting	r	j	Δd_t	100 Δ bbox
COCO B/16	18	12	0.029	-19.53
			[-0.096,0.148]	[-38.12,-1.32]
COCO B/32	30	25	0.094	-14.04
			[0.036,0.148]	[-27.70,-1.74]
VOC B/16	5	3	-0.021	-16.87
			[-0.170,0.182]	[-26.11,-7.00]
VOC B/32	2	1	-0.078	-28.17
			[-0.115,-0.042]	[-78.57,22.22]

Repairing specificity can reduce spatial overlap. Table 5 considers only repaired failures; Table 3 uses all eligible images. Mean bbox precision falls by 14.04–28.17 percentage points among repairs. Its interval is below zero in three settings, but spans zero for VOC B/32. Mean target-score changes are positive on COCO and negative on VOC; their intervals span zero for both B/16 settings. Thus the direction of a mean change is not always resolved by its interval. VOC estimates summarize just five and two repairs.

A repaired margin also need not retain the other checks. On COCO B/32, 30 failures are repaired, but only 25 retain overlap, mean-foil contrast, and positive target contribution. This gap between repairs (r) and repairs passing all checks (j) occurs in every setting. Even the oracle finds only 12, 26, 3, and 2 feasible repairs passing all checks, in table order. Among all baseline passes on COCO B/32, four instead become worst-foil failures; these regressions lie outside the fixed failure set.

Table 6: Nonzero-threshold sensitivity on fixed baseline A . Entries are percentages of screen-passers A with $P < -\delta$, where $\delta \geq 0$ is the required failure magnitude in the normalized margin. Thresholds are evaluated post hoc in logit units.

Setting	$\delta = 0$	.01	.10	.50
COCO B/16	64.37	63.94	59.49	36.05
COCO B/32	64.78	64.38	60.22	35.19
VOC B/16	42.27	41.44	36.23	13.73
VOC B/32	41.08	40.26	34.39	11.90

The failures are not confined to margins near zero. Table 6 tightens the test from $P < 0$ to $P < -\delta$. At $\delta = .10$, 59.49–60.22% of COCO screen-passers and 34.39–36.23% of VOC screen-passers still fail. Rates decline further at .50. Thus many failures survive a stricter margin requirement, although their prevalence depends on the chosen threshold.

Implications for region selection. The comparisons point to two distinct changes. Table 2 identifies images that need a better candidate region (C_0) and images that already have a passing region outside the budget (C_1). Figure 2 shows how sharply the budget narrows the available choices, and Table 3 shows the effect of relaxing it. Changing the ranking rule alone can recover only the additional feasible opportunities in C_2 . Improving repair therefore requires attention to which regions are proposed and how much target-drop loss is allowed.

Reproducibility and uncertainty. Main results use a fully recorded rerun of the earlier seed rule, not independent base-seed repetitions. Checks reconstruct regions, scores, choices, and failures from saved masks, responses, and bbox records. Released records include summaries, checks, configuration, and hashes; masks and full responses remain local. Table 4 uses earlier-run records. Selected regions changed in three images.

We obtain confidence intervals from 10,000 image-level bootstrap samples, using the 2.5th and 97.5th percentiles without multiplicity correction or coverage of base-seed variability. Paired effects are differenced within each image. Foil comparisons use fixed A or A^+ ; repair-cost intervals use only actual repairs. We use exact q_i values, checked against 100 seeded random draws.

5 CONCLUSION

A region can overlap the target and beat the average competing class while contributing more to a single competitor. This occurs frequently for CCI-top1, even after annotated competitors are removed. Checking every candidate shows why reranking repairs so few failures: many images have no passing alternative, and most remaining alternatives exceed the target-drop tolerance. Relaxing that tolerance makes more repairs possible but also changes target response and overlap. Worst-foil comparisons therefore add information that overlap and mean contrast alone do not provide when evaluating frozen CLIP region explanations.

REFERENCES

- [1] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, G. Krueger, and I. Sutskever, “Learning transferable visual models from natural language supervision,” in *Proceedings of the 38th International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, vol. 139. PMLR, 2021, pp. 8748–8763. [Online]. Available: <https://proceedings.mlr.press/v139/radford21a.html>.
- [2] A. Agarwal, S. Karanam, and V. Gandhi, “Concept regions matter: Benchmarking CLIP with a new cluster-importance approach,” *arXiv preprint arXiv:2511.12978*, 2025. [Online]. Available: <https://arxiv.org/abs/2511.12978>.
- [3] Y. Wang and X. Wang, “‘Why not other classes?’: Towards class-contrastive back-propagation explanations,” in *Advances in Neural Information Processing Systems*, vol. 35, 2022. [Online]. Available: <https://doi.org/10.52202/068431-0660>.
- [4] C. Lin, H. Chen, C. Kim, and S.-I. Lee, “Contrastive corpus attribution for explaining representations,” in *International Conference on Learning Representations (ICLR)*, 2023. [Online]. Available: <https://arxiv.org/abs/2210.00107>.
- [5] D. Williamson, Y. Ji, and M. Dwyer, “CASE: Contrastive activation for saliency estimation,” *arXiv preprint arXiv:2506.07327*, 2025. [Online]. Available: <https://arxiv.org/abs/2506.07327>.
- [6] Y. Wang, T. G. J. Rudner, and A. G. Wilson, “Visual explanations of image-text representations via multi-modal information bottleneck attribution,” in *Advances in Neural Information Processing Systems*, vol. 36, 2023, pp. 16 009–16 027. [Online]. Available: <https://doi.org/10.52202/075280-0704>.
- [7] Y. Li, H. Wang, Y. Duan, J. Zhang, and X. Li, “A closer look at the explainability of contrastive language-image pre-training,” *Pattern Recognition*, vol. 162, p. 111409, 2025. [Online]. Available: <https://doi.org/10.1016/j.patcog.2025.111409>.
- [8] M. Dreyer, L. Hufe, J. Berend, T. Wiegand, S. Lapuschkin, and W. Samek, “From what to how: Attributing CLIP’s latent components reveals unexpected semantic reliance,” *arXiv preprint arXiv:2505.20229*, 2025. [Online]. Available: <https://arxiv.org/abs/2505.20229>.
- [9] F. Li, Z. Meng, C. Tan, and Y. Yin, “Class-disentangled attribution with gradient guidance for CLIP,” *Journal of King Saud University Computer and Information Sciences*, vol. 38, p. 376, 2026. [Online]. Available: <https://doi.org/10.1007/s44443-026-00779-3>.
- [10] M. Yuksekgonul, F. Bianchi, P. Kalluri, D. Jurafsky, and J. Zou, “When and why vision-language models behave like bags-of-words, and what to do about it?” in *International Conference on Learning Representations (ICLR)*, 2023. [Online]. Available: <https://openreview.net/forum?id=KRLUvxh8uaX>.
- [11] T.-Y. Lin, M. Maire, S. Belongie, J. Hays, P. Perona, D. Ramanan, P. Dollar, and C. L. Zitnick, “Microsoft COCO: Common objects in context,” in *Computer Vision—ECCV 2014*, ser. Lecture Notes in Computer Science, vol. 8693. Springer, 2014, pp. 740–755. [Online]. Available: https://doi.org/10.1007/978-3-319-10602-1_48.
- [12] M. Everingham, L. Van Gool, C. K. I. Williams, J. Winn, and A. Zisserman, “The PASCAL visual object classes (VOC) challenge,” *International Journal of Computer Vision*, vol. 88, no. 2, pp. 303–338, 2010. [Online]. Available: <https://doi.org/10.1007/s11263-009-0275-4>.
- [13] T. Gomez, T. Freour, and H. Mouchere, “Metrics for saliency map evaluation of deep learning explanation methods,” in *Pattern Recognition and Artificial Intelligence: Third International Conference, ICPRAI 2022*, ser. Lecture Notes in Computer Science. Springer, 2022, pp. 84–95. [Online]. Available: https://doi.org/10.1007/978-3-031-09037-0_8.

COMPLIANCE WITH ETHICAL STANDARDS

This work analyses publicly released COCO and PASCAL VOC images and annotations. No participants were recruited and no new human or animal experiments were conducted.

ACKNOWLEDGMENTS

This work was supported by the Young Scientific and Technological Talent Support Program under the Taizhou Fengcheng Talent Plan.

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.